



# Using R's Caret Package for Machine Learning

Earl F Glynn

Kansas City R Users Group

9 September 2017


<https://github.com/EarlGlynn/kc-r-users-caret-2017>

# Using R's Caret Package for Machine Learning

- Regressions Models vs Machine Learning
- Caret Package Overview
  - Visualize Data
  - Preprocess / Transform Data
  - Partition Data into Train and Test Sets
  - Train and Tune Model
  - Evaluate Model Performance
  - Estimate Variable Importance
  - Parallel Processing
- Caret Machine Learning: Forensic Glass Dataset
- Toy Shiny App to Compare Machine Learning Models
- Summary

# Regression Models vs. Practical Machine Learning

**coursera**



Johns Hopkins University  
Regression Models

**Focus: Interpretability**




Johns Hopkins University  
Practical Machine Learning

**Focus: Accurate Predictions**

<https://www.coursera.org/specializations/jhu-data-science>

# Caret Package: Classification And Regression Training



2013

<http://appliedpredictivemodeling.com/>  
<http://appliedpredictivemodeling.com/blog/>



Max Kuhn

topepo

<https://github.com/topepo/caret>



Hadley Wickham

@hadleywickham

Following

Super excited to announce that Max Kuhn is joining my team at @rstudio:

[appliedpredictivemodeling.com/blog/2016/11/2...](http://appliedpredictivemodeling.com/blog/2016/11/2...) 🎉🎉🎉 #rstats



#### Working at RStudio

I've joined Hadley's team at RStudio. Unsurprisingly, I'll be working on some modeling related R packages and infrastructure. It is very exciting and I'm looking forward to learning a lot and...  
[appliedpredictivemodeling.com](http://appliedpredictivemodeling.com)

9:23 PM - 28 Nov 2016

# Caret Package Features

- Provides uniform interface to machine learning models
- Streamlines model training and tuning
- Standardizes common tasks
- Uses dozens of R packages
- Provides parallel processing

Building Predictive Models in R Using the caret Package

<https://www.jstatsoft.org/article/view/v028i05/v28i05.pdf>

Max Kuhn, Pfizer Global R&D

*Journal of Statistical Software*, Nov. 2008

# Caret Package Models

## 6 Available Models

The models below are available in `train`. The code behind these protocols can be obtained using the function `getModelInfo` or by going to the [github repository](#).

Show  entries

Search:

| Model  | <i>method</i> Value | Type                       | Libraries          |
|--|---------------------|----------------------------|--------------------|
| AdaBoost Classification Trees                          | adaboost            | Classification             | fastAdaboost       |
| AdaBoost.M1  | AdaBoost.M1         | Classification             | adabag, plyr       |
| Adaptive Mixture Discriminant Analysis                 | amdai               | Classification             | adaptDA            |
| Adaptive-Network-Based Fuzzy Inference System          | ANFIS               | Regression                 | frbs               |
| Adjacent Categories Probability Model for Ordinal Data | vglmAdjCat          | Classification             | VGAM               |
| Bagged AdaBoost  | AdaBag              | Classification             | adabag, plyr       |
| Bagged CART  | treebag             | Classification, Regression | ipred, plyr, e1071 |

<http://topepo.github.io/caret/modelList.html>

Files: **caret-overview.Rmd** and **caret-overview.html**

# Forensic Glass Dataset



<https://sha.org/bottle/links.htm>

*fgl* dataset in MASS package

- RI = refractive index
- Percentages by weight of oxides: Na, Mg, Al, Si, K, Ca, Ba, Fe
- type
  - window float glass (WinF: 70),
  - window non-float glass (WinNF: 76),
  - vehicle window glass (Veh: 17),
  - containers (Con: 13),
  - tableware (Tabl: 9)
  - vehicle headlamps (Head: 29).

| RI    | Na    | Mg   | Al   | Si    | K    | Ca   | Ba | Fe   | type |
|-------|-------|------|------|-------|------|------|----|------|------|
| 3.01  | 13.64 | 4.49 | 1.10 | 71.78 | 0.06 | 8.75 | 0  | 0.00 | WinF |
| -0.39 | 13.89 | 3.60 | 1.36 | 72.73 | 0.48 | 7.83 | 0  | 0.00 | WinF |


Also available through [UCI Repository](#).

Discussed in book [Data Mining and Business Analytics with R](#).

# Caret: Parallel Processing

On a PC ...

```
# Setup parallel processing
# Let's use 6 cores
library(doParallel)
rCluster <- makePSOCKcluster(6)
registerDoParallel(rCluster)
...
stopCluster(rCluster)
```






# Caret: Visualize Data

## Pairs Plots

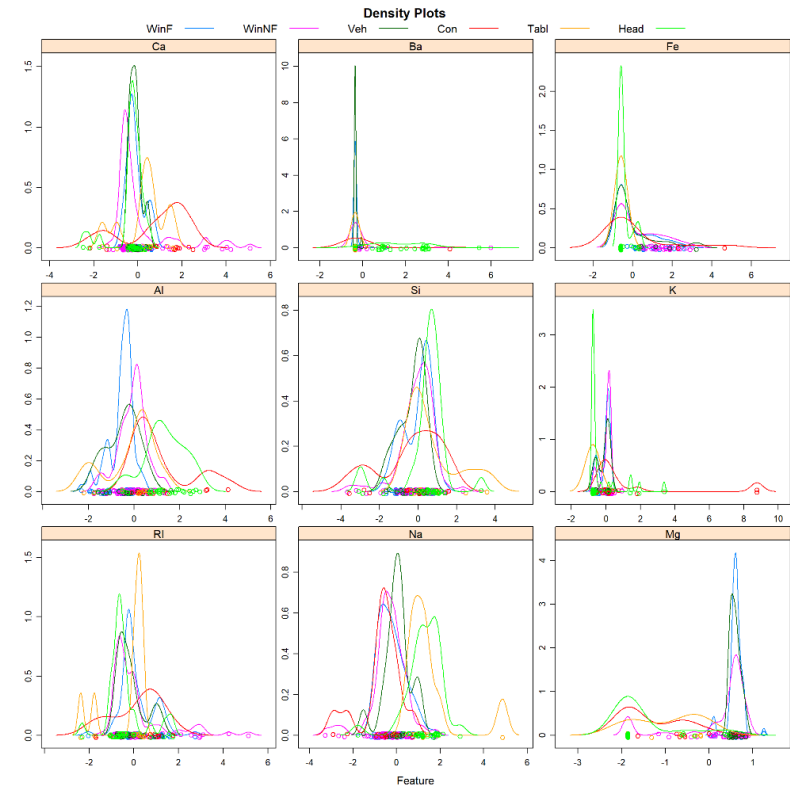
```
featurePlot(centeredScaled, type, "pairs", auto.key=list(columns=6), main="Pairs Plots")
```



## Density Plots

See [Density estimation](#).

```
featurePlot(centeredScaled, type, "density",  
  scales=list(x=list(relation="free"),  
    y=list(relation="free")),  
  auto.key=list(columns=6),  
  main="Density Plots")
```



Files: **Forensic-Glass-Exploratory.Rmd** or **Forensic-Glass-Exploratory.html**

Data Visualization with the Caret R package

<https://machinelearningmastery.com/data-visualization-with-the-caret-r-package/>

# Caret: Partition Data into Train & Test

Use function **createDataPartition** to create splits of the data

```
library(MASS)
```

```
library(caret)
```

## Forensic Glass Data

```
rawData <- fgl  
dim(rawData)
```

```
[1] 214 10
```

## Define train and test datasets

```
set.seed(71)  
  
trainSetIndices <- createDataPartition(rawData$type, p=0.70, list=FALSE)  
  
trainSet <- rawData[ trainSetIndices, ]  
testSet  <- rawData[ -trainSetIndices, ]
```

# Caret: Partition Data into Train & Test

Use function **createDataPartition** to create splits of the data.

<http://topepo.github.io/caret/data-splitting.html>

Approach Using Forensic Glass Dataset

214 glass samples each with 9 predictors

Split original dataset: 70% training, 30% final test

153 samples  
training

Use for training and tuning

61 samples  
final test

Use once and only once for  
final test of model. Why?

# Caret: Preprocess / Transform Data

- Impute Missing Values
- Create Dummy Variables
- Remove Zero- / Near Zero-Variance Predictors
- Remove Correlated Predictors
- Remove Linear Dependencies
- Center / Scale / Standardize Data
- Transformations (BoxCox, YeoJohnson)

<http://topepo.github.io/caret/pre-processing.html>

# Caret: Train Model

## Linear Discriminant Analysis

```
set.seed(29)
CVfolds <- 5 # 5-fold cross validation (not enough data for 10 fold here)
CVrepeats <- 10 # repeat 10 times

# Used createMultiFolds to study
indexFolds <- createMultiFolds(trainSet$type, CVfolds, CVrepeats) # for repeated CV


trainControlParms <- trainControl(method = "repeatedcv", # repeated cross validation
                                   number = CVfolds,
                                   repeats = CVrepeats,
                                   index = indexFolds,
                                   classProbs = TRUE, # Estimate class probabilities
                                   summaryFunction = defaultSummary)

fit <- train(type ~ ., data=trainSet,
             preProcess = c("center", "scale"),
             method = "lda",
             metric = "Kappa",
             trControl = trainControlParms)
```

Files: **Forensic-Glass-caret-rf.Rmd** and **Forensic-Glass-caret-rf.html**


# Caret: Train Model

## K-Fold Cross-Validation




# Underfit vs Overfit


## Bias-Variance Tradeoff



High bias  
(underfit)



"Just right"



High variance  
(overfit)

# Caret: Tune Model

Often default tuning is adequate, but when not ...

Hyperparameters

J48

```
set.seed(29)
CVfolds <- 5 # 5-fold cross validation (not enough data for 10 fold here)
CVrepeats <- 10 # repeat 10 times

TUNEgrid <- expand.grid(C = 1:4 * 0.1,
                       M = 1:5)

trainControlParms <- trainControl(method = "repeatedcv", # repeated cross validation
                                   number = CVfolds,
                                   repeats = CVrepeats,
                                   classProbs = TRUE,      # Estimate class probabilities
                                   summaryFunction = defaultSummary)

fit <- train(type ~ ., data=trainSet,
             method = "J48",
             metric = "Kappa", # helps with imbalance
             tuneGrid = TUNEgrid, # expanded range
             trControl = trainControlParms)
```

| C   | M |
|-----|---|
| 0.1 | 1 |
| 0.1 | 2 |
| 0.1 | 3 |
| 0.1 | 4 |
| 0.1 | 5 |
| 0.2 | 1 |
| 0.2 | 2 |
| 0.2 | 3 |
| 0.2 | 4 |
| 0.2 | 5 |
| 0.3 | 1 |
| 0.3 | 2 |
| 0.3 | 3 |
| 0.3 | 4 |
| 0.3 | 5 |
| 0.4 | 1 |
| 0.4 | 2 |
| 0.4 | 3 |
| 0.4 | 4 |
| 0.4 | 5 |

Files: **Forensic-Glass-caret-J48.Rmd** and **Forensic-Glass-caret-J48f.html**

Tune Machine Learning Algorithms in R (random forest case study)  
<https://machinelearningmastery.com/tune-machine-learning-algorithms-in-r/>



# HYPERPARAMETER TUNING

Finding the hyperparameter values of a learning algorithm that produce the best model.

ChrisAlbon

## PARAMETERS VS. HYPERPARAMETERS

Parameters are learned through the training procedure. For example, the weights of a neural network.

Hyperparameters are set before training starts, can be tuned through grid search or related methods.

ChrisAlbon

# Caret: Evaluate Model Performance

## Classification Metrics

- Accuracy
- Area Under Curve (AUC) [2 classes]
- Kappa (better than Accuracy when groups not balanced)
- LogLoss

## Regression Metrics

- RMSE (Root Mean Squared Error)
- $R^2$

# Caret: Evaluate Model Performance

## ACCURACY

$$A_{cc} = \frac{1}{n} \sum 1(\hat{y}_i = y_i)$$

*n* ← number of observations  
Indicator function  
Predicted  $y$   
True  $y$

A common metric in classification. Fails when we have highly imbalanced classes. In those cases F1 is more appropriate.

ChrisAlbon

## Kappa

For classification models:

- **overall accuracy** can be used, but this may be problematic when the classes are not balanced.
- the **Kappa statistic** takes into account the expected error rate:

$$\kappa = \frac{O - E}{1 - E}$$

where  $O$  is the observed accuracy and  $E$  is the expected accuracy under chance agreement (`psych::cohen.kappa`, `vcd::Kappa`, ...)

[https://www.r-project.org/conferences/user-2013/Tutorials/kuhn/user\\_caret\\_2up.pdf](https://www.r-project.org/conferences/user-2013/Tutorials/kuhn/user_caret_2up.pdf)

## CONFUSION MATRIX

Confusion matrices visualize the accuracy of a classifier by comparing the true and predicted classes. Off diagonal squares are incorrect predictions.

|            |   | Predicted Class |   |   |
|------------|---|-----------------|---|---|
|            |   | A               | B | C |
| True Class | A |                 |   |   |
|            | B |                 |   |   |
|            | C |                 |   |   |

ChrisAlbon

@ChrisAlbon

MachineLearningFlashcards.com

# Caret: Evaluate Model Performance

- **Sensitivity:** given that a result is truly an event, what is the probability that the model will predict an event results?
- **Specificity:** given that a result is truly not an event, what is the probability that the model will predict a negative results?

[https://www.r-project.org/conferences/useR-2013/Tutorials/kuhn/user\\_caret\\_2up.pdf](https://www.r-project.org/conferences/useR-2013/Tutorials/kuhn/user_caret_2up.pdf)

# Caret: Evaluate Model Performance

## Train: Overly Optimistic Results for Generalization

### Results on Train Set (In Sample)

Overly optimistic results for generalization

```
options(width=120)
InSample <- predict(fit, newdata=trainSet)
InSampleConfusion <- confusionMatrix(trainSet$type, InSample)
print(InSampleConfusion)
```

#### Confusion Matrix and Statistics

|            | Reference |       |     |     |      |      |
|------------|-----------|-------|-----|-----|------|------|
| Prediction | WinF      | WinNF | Veh | Con | Tabl | Head |
| WinF       | 49        | 0     | 0   | 0   | 0    | 0    |
| WinNF      | 0         | 54    | 0   | 0   | 0    | 0    |
| Veh        | 0         | 0     | 12  | 0   | 0    | 0    |
| Con        | 0         | 0     | 0   | 10  | 0    | 0    |
| Tabl       | 0         | 0     | 0   | 0   | 7    | 0    |
| Head       | 0         | 0     | 0   | 0   | 0    | 21   |

#### Overall Statistics

Accuracy : 1

95% CI : (0.9762, 1)

No Information Rate : 0.3529

P-Value [Acc > NIR] : < 2.2e-16

Kappa : 1


Mcnemar's Test P-Value : NA

#### Statistics by Class:

|             | Class: WinF | Class: WinNF | Class: Veh | Class: Con | Class: Tabl | Class: Head |
|-------------|-------------|--------------|------------|------------|-------------|-------------|
| Sensitivity | 1.0000      | 1.0000       | 1.00000    | 1.00000    | 1.00000     | 1.0000      |
| Specificity | 1.0000      | 1.0000       | 1.00000    | 1.00000    | 1.00000     | 1.0000      |

## OVERFITTING

Overfitting occurs when a model starts to memorize the aspects of the training set and in turn loses the ability to generalize



ChrisAlbon

@ChrisAlbon  
MachineLearningFlashcards.com

Files: **Forensic-Glass-caret-rf.Rmd** and **Forensic-Glass-caret-rf.html**

# Caret: Evaluate Model Performance

## Test: More Realistic Results on Predictions with New Data

### Results on Test Set (Out of Sample)

More realistic results on predictions with new data

```
options(width=120)
OutOfSample <- predict(fit, newdata=testSet)
confusion <- confusionMatrix(testSet$type, OutOfSample)
print(confusion)
```

#### Confusion Matrix and Statistics

|            | Reference |       |     |     |      |      |
|------------|-----------|-------|-----|-----|------|------|
| Prediction | WinF      | WinNF | Veh | Con | Tabl | Head |
| WinF       | 18        | 3     | 0   | 0   | 0    | 0    |
| WinNF      | 4         | 15    | 0   | 2   | 1    | 0    |
| Veh        | 1         | 1     | 3   | 0   | 0    | 0    |
| Con        | 0         | 0     | 0   | 2   | 0    | 1    |
| Tabl       | 0         | 0     | 0   | 0   | 2    | 0    |
| Head       | 0         | 0     | 0   | 1   | 0    | 7    |

#### Overall Statistics

Accuracy : 0.7705

95% CI : (0.645, 0.8685)

No Information Rate : 0.377

P-Value [Acc > NIR] : 4.448e-10

Kappa : 0.686

Mcnemar's Test P-Value : NA

#### Statistics by Class:

|             | Class: WinF | Class: WinNF | Class: Veh | Class: Con | Class: Tabl | Class: Head |
|-------------|-------------|--------------|------------|------------|-------------|-------------|
| Sensitivity | 0.7826      | 0.7895       | 1.00000    | 0.40000    | 0.66667     | 0.8750      |
| Specificity | 0.9211      | 0.8333       | 0.96552    | 0.98214    | 1.00000     | 0.9811      |

Files: **Forensic-Glass-caret-rf.Rmd** and **Forensic-Glass-caret-rf.html**


# Caret: Estimate Variable Importance

## Random Forest: Single Prediction Model

Variable Importance

See ?varImp

```
plot( varImp(fit), main="Variable Importance" )
```



Files: **Forensic-Glass-caret-rf.Rmd** and **Forensic-Glass-caret-rf.html**


# Caret: Estimate Variable Importance

## J48: Separate Model for Each Type

Variable Importance

See ?varImp

```
plot( varImp(fit), main="Variable Importance" )
```



Files: **Forensic-Glass-caret-J48.Rmd** and **Forensic-Glass-caret-J48.html**



# Caret Machine Learning

## Which Machine Learning Algorithms to use?

### Theorem (No Free Lunch)

*In the absence of any knowledge about the prediction problem, no model can be said to be uniformly better than any other*

Given this, it makes sense to use a variety of different models to find one that best fits the data

# Caret Machine Learning

## Which Machine Learning Algorithms to use?

### Algorithms

It is important to have a good mix of algorithm representations (lines, trees, instances, etc.) as well as algorithms for learning those representations.

A good rule of thumb I use is “a few of each”, for example in the case of binary classification:

- **Linear methods:** Linear Discriminant Analysis and Logistic Regression.
- **Non-Linear methods:** Neural Network, SVM, kNN and Naive Bayes
- **Trees and Rules:** CART, J48 and PART
- **Ensembles of Trees:** C5.0, Bagged CART, Random Forest and Stochastic Gradient Boosting

You want some low complexity easy to interpret methods in there (like LDA and kNN) in case they do well, you can adopt them. You also want some sophisticated methods in there (like random forest) to see if the problem can even be learned and to start building up expectations of accuracy.

How many algorithms? At least 10-to-20 different algorithms.

# Caret Machine Learning Examples

| Group              | Algorithms  | Caret Model      | FILE             |
|--------------------|---|------------------|------------------|
| Linear Methods     | Linear Discriminant Analysis  | lda              | lda              |
|                    | Linear Discriminant Analysis with YeoJohnson preprocessing                            | lda w/YeoJohnson | lda-YeoJohnson   |
|                    | LASSO, Ridge, and Elastic Net   | glmnet           | glmnet           |
|                    | LASSO, Ridge, and Elastic Net with Synthetic Minority Over-Sampling Technique (SMOTE) | glmnet w/SMOTE   | glmnet-SMOTE     |
| Non-Linear Methods | Neural Network  | nnet             | nnet             |
|                    | Support Vector Machine with Radial Basis Function Kernel                              | svmRadial        | svmRadial        |
|                    | Naïve Bayes   | nb               | nb               |
|                    | Naïve Bayes with Independent Component Analysis (ICA)                                 | nb w/ICA         | nb-ica-transform |
|                    | k Nearest Neighbors   | knn              | knn              |
| Trees and Rules    | J48   | J48              | J48              |
|                    | Classification and Regression Trees   | rpart            | rpart            |
| Ensembles of Trees | C5.0  | C5.0             | C50              |
|                    | Random Forest   | rf               | rf               |
|                    | Random Forest with SMOTE  | rf w/SMOTE       | rf-SMOTE         |

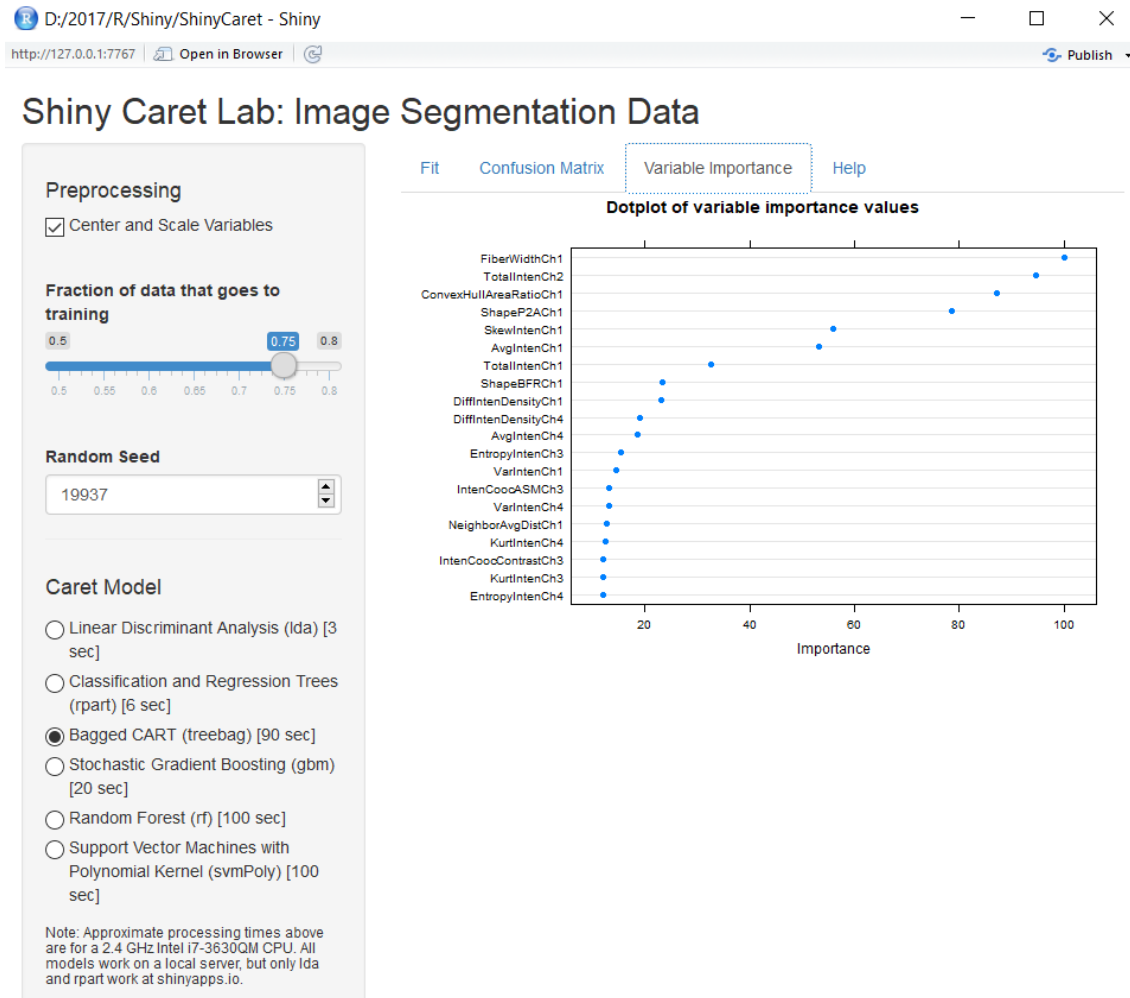
Files: **Forensic-Glass-caret-FILE.Rmd** and **Forensic-Glass-caret-FILE.html**

# Summary of Results From Variety of Machine Learning Methods

Hold Out Test Dataset Results

| Group              | Model            | Accuracy     | Kappa        |
|--------------------|------------------|--------------|--------------|
| Linear Methods     | lda              | 0.541        | 0.369        |
|                    | lda w/YeoJohnson | 0.623        | 0.486        |
|                    | glmnet           | 0.639        | 0.498        |
|                    | glmnet w/SMOTE   | 0.672        | 0.547        |
| Non-Linear Methods | nnet             | 0.672        | 0.555        |
|                    | svmRadial        | 0.639        | 0.486        |
|                    | nb               | 0.607        | 0.459        |
|                    | nb w/ICA         | 0.508        | 0.314        |
|                    | knn              | 0.689        | 0.582        |
| Trees and Rules    | J48              | 0.672        | 0.560        |
|                    | rpart            | 0.574        | 0.373        |
| Ensembles of Trees | C5.0             | 0.689        | 0.573        |
|                    | rf               | <b>0.771</b> | <b>0.686</b> |
|                    | rf w/SMOTE       | 0.754        | 0.666        |

# Toy Shiny App To Compare Models



Open ShinyCaret Folder as project in RStudio. Open file server.R. Select Run App.

# Machine Learning Algorithms Using R's Caret Package

## Summary


- Caret provides uniform approach to using many classification and regression algorithms
- Many machine algorithms can be explored quickly using Caret.
- Caret provides many useful tools for machine learning

# Question

*Do We Need Hundreds of Classifiers to Solve Real World Classification Problems?*

<http://jmlr.org/papers/volume15/delgado14a/delgado14a.pdf>

# Forensic Glass Dataset: PCA



WinF WinNE Veh Con Tabl Head

Files: **Forensic-Glass-PCA.Rmd** and **Forensic-Glass-PCA.html**